

IC3PO Finite Convergence Check

Ian Dardik

February 3, 2022

1 Introduction

IC3PO proposed an automatic method for performing “finite convergence checks” [?]. In this document I will describe the problem, the proposed algorithm, as well as issues that I discovered. Throughout this document I will adopt the conventions as used in [?], but I have chosen different variable names for convenience and readability. There are several places in this document that ought to have proofs, but each one is very simple and I can include them if anyone is interested.

2 Preliminaries

Let P be either a property or a transition system. Then the *template* of P is $P(S_1, \dots, S_n)$ where $n \in \mathbb{N}$ and each S_i is a sort. A finite instance of P is denoted as $P(|S_1|, \dots, |S_n|)$, where each $|S_i| \in \mathbb{N}$ denotes the size of sort S_i . For example, if we have a property *Init* parameterized by two sorts S_1 and S_2 , then $\text{Init}(S_1, S_2)$ is the template of *Init* while $\text{Init}(2, 3)$ is a finite instance.

3 Problem Statement

Given a transition system $T = (\text{Init}, \text{Trans})$ with template $T(S_1, \dots, S_n)$ and a property P with template (S_1, \dots, S_n) , we want to determine whether $P(S_1, \dots, S_n)$ is an inductive invariant for $T(S_1, \dots, S_n)$. $P(S_1, \dots, S_n)$ is said to be an inductive invariant for $T(S_1, \dots, S_n)$ iff $P(|S_1|, \dots, |S_n|)$ is an inductive invariant for $T(|S_1|, \dots, |S_n|)$ for all $\{|S_1|, \dots, |S_n|\} \in \mathbb{N}^n$.

4 Proposed Solution

Given the problem statement outlined in the previous section, [?] proposes the following algorithm: Manually identify a vector $V \in \mathbb{N}^n$ which assigns *basesizes* to each of the n sorts. V must be chosen such that $T(v_1, \dots, v_n)$ exhibits non-trivial behavior; “non-trivial behavior” is glossed over in the IC3PO paper but is intuitively sort sizes large enough to see some complex behaviors in the protocol T . Then P is an inductive invariant if the following checks hold for $1 \leq i \leq n$:

1. $\text{Init}(v_1, \dots, v_i + 1, \dots, v_n) \rightarrow P(v_1, \dots, v_i + 1, \dots, v_n)$
2. $P(v_1, \dots, v_i + 1, \dots, v_n) \wedge \text{Trans}(v_1, \dots, v_i + 1, \dots, v_n) \rightarrow P'(v_1, \dots, v_i + 1, \dots, v_n)$

5 Counterexample

There is a simple counterexample to this proposed solution. We begin by defining the properties Q and R and then proceed to describe the counterexample.

5.1 The Q and R Properties

Suppose that $T = (Init, Trans)$ is a transition system parameterized by a single sort S , and has at least one reachable state outside of $Init$. We will make use of two properties:

1. Q , which is any inductive invariant for $T(S)$
2. $R(m) := (|S| = m) \rightarrow Init$

Because T has at least one reachable state outside of $Init$, it is clear that $Init$ is neither inductive nor invariant. It follows that there is no value of $m \in \mathbb{N}$ such that $Q \wedge R(m)$ is an inductive invariant for $T(S)$.

5.2 $R(1)$ Counterexample

The proposed solution in section 4 will incorrectly decide that $Q \wedge R(1)$ is an inductive invariant. To see why, consider an arbitrary basesize vector $V = \{v\}$ where $v \in \mathbb{N}$. Then IC3PO will check the following for finite convergence:

1. $Init(v+1) \rightarrow (Q \wedge R(1))(v+1)$
2. $(Q \wedge R(1))(v+1) \wedge Trans(v+1) \rightarrow (Q \wedge R(1))'(v+1)$

Notice that if $v > 0$, then

$$\begin{aligned}
 & (Q \wedge R(1))(v+1) \\
 \iff & Q(v+1) \wedge (v+1 = 1 \rightarrow Init) \\
 \iff & Q(v+1) \wedge (false \rightarrow Init) \\
 \iff & Q(v+1)
 \end{aligned}$$

Thus, for any basesize larger than 0, the finite convergence check reduces to checking whether $Q(v+1)$ is an inductive invariant. Because Q is an inductive invariant, the algorithm will incorrectly decide that $Q \wedge R(1)$ is an inductive invariant.

5.3 $R(1)$ Counterexample in IC3PO

The $R(1)$ counterexample is easily realized in IC3PO with the following Ivy program:

```

#lang ivy1.7
type num
relation val(X:num)
after init {
  val(X) := false;
}
action t(x:num) = {
  val(x) := true;
}
export t
invariant [Q] val(X) | ~val(X)
invariant [R(1)] (forall X:num, Y. X = Y) -> (forall Z. ~val(Z))

```

In this Ivy program there is a single sort *num*. Notice that in this program:

$$\begin{aligned}
& Q \\
& \iff \forall X : \text{val}(X) \vee \neg \text{val}(X) \\
& \iff \text{true}
\end{aligned}$$

which is clearly an inductive invariant, and

$$\begin{aligned}
& R(1) \\
& \iff (\forall X, Y : X = Y) \rightarrow (\forall Z, \neg \text{val}(Z)) \\
& \iff (|num| = 1) \rightarrow \text{Init}
\end{aligned}$$

Running IC3PO on this program with $|num| = 0$ yields a property violation as expected. Running with $|num| = 1$ yields a property violation too; this is expected because IC3PO finds the property violation while it tries to synthesize an invariant (not during the finite convergence check). For any values of $|num| > 1$ however, IC3PO incorrectly posits that the conjunction of Q and $R(1)$ is an inductive invariant. Since this example is in EPR, I was able to run *ivy_check* (which does not use finite interpretations of the sorts) to produce a counterexample and verify that $Q \wedge R(1)$ is not an inductive invariant.

5.4 $R(m)$ Counterexamples

This counterexample generalizes to higher values of m , though I have not tested it on IC3PO yet (because I don't know how to test for cardinality of a sort in Ivy). Imagine the following Ivy program:

```

#lang ivy1.7
type num
relation val(X:num)
after init {
  val(X) := false;
}
action t(x:num) = {
  val(x) := true;
}
export t
invariant [Q] val(X) | ~val(X)
invariant [R(100)] (|num|=100) -> (forall Z. ~val(Z))

```

In this case, the conjunction of the invariants ($Q \wedge R(100)$) is still clearly not an inductive invariant. However, users of IC3PO will most likely choose basesizes of small integers such as 2 or 3; checking for convergence using the method in 4 will incorrectly decide that the conjunction of the invariants *is* an inductive invariant. In this case, the user would need to choose a basesize of 99 or 100 for the method to work properly.

6 Discussion

Since m can be arbitrarily chosen to create $R(m)$, *no* finite number of checks on finite instances are sufficient to prove that a property is an inductive invariant. However, it is possible that there is some restricted class of properties for which a finite number of checks does suffice.

7 Conclusion

A counterexample to IC3PO's finite convergence detection algorithm is presented. It is possible that restricting the universe of allowed inductive invariants could make the method (or a similar one) work, and this would be a very interesting research direction.