A Cutoff Rule For A Special Class Of Parameterized Distributed Protocols

Ian Dardik

March 7, 2022

1 Introduction

In this note, we consider the verification problem of a transition system $T = (I, \Delta)$ where I is the initial constraint, Δ is the transition relation, and the system is parameterized by a single sort P of identical elements (We make the notion of "identical" precise in Assumption 1 below). We assume that we are given a candidate inductive invariant Φ which implies our key safety property. Φ is restricted to be in Prenex Normal Form (PNF) with only universal quantifiers, while Δ is restricted to be in PNF with only existential quantifiers. We adopt the convention of [2] where $T(P)$ is the template of T, and $T(|P|)$ is a finite instantiation.

In this note, we will build several lemmas that lead to an interesting result: let m be the number of variables that Φ quantifies over and n be the number of variables that Δ quantifies over, then if $\Phi(m+n)$ is an inductive invariant, $\Phi(k)$ is also an inductive invariant for all $k > m+n$. We will refer to this as the M-N Theorem in this note. This result is useful because it reduces the vefification problem on T to model checking a finite number of instances $T(1), T(2), ..., T(m+n)$. Essentially, $m+n$ is a cutoff instance size for proving that our inductive invariant holds.

Note: I think it is likely that if $\Phi(m+n)$ is an inductive invariant, then it is *also* the case for $\Phi(k)$ for all $k < m+n$, but I left this out of this note for the time being to focus on the $k > m+n$ case.

2 Preliminaries

In this section we cover several preliminary items that we use to prove the M-N Theorem.

2.1 Without Loss Of Generality

We will assume that the parameter $P = \{1, 2, ..., |P|\}$. This assumption comes without loss of generality because each member of P is assumed to be identical.

2.2 Assumptions

This section contains the list of assumptions for the transition system we work with. In other words, these assumptions are the requirements for the M-N Theorem to hold.

Assumption 1 (P Has Identical Elements)**.** Let f be a ground formula and let $\pi : P \to P$ be a bijective function, i.e. a permutation on P. Then we assume:

$$f \leftrightarrow \pi(f)$$

2.3 Definitions

Definition 1 (States)**.** Let $k \in \mathbb{N}$, then:

$$\text{States}(k) := \{s \mid s \text{ is a state of } T(k)\}$$

In this note we consider a state $s \in \text{States}(k)$ to be a quantifier-free formula: a conjunction of constraints that describe a single state in $T(k)$.

Definition 2 (Satisfaction)**.** Let f and g be formulas in First Order Logic. Then we say $f \models g$ iff $f \to g$. Alternatively, f satisfies g iff f is stronger than g.

Example 1. Consider the transition system $T(P)$ with two state variables, $x \in (P \to \mathbb{N})$ and $y \in \mathbb{Z}$. Let $|P| = 2$, then $P = \{1, 2\}$. Let $s := (x[1] = 6 \wedge x[2] = 0 \wedge y = -22)$ be a state in the transition system. Let $F := \forall p, q \in P, x[p] \neq x[q]$ and $f := (x[1] \neq x[2])$. Then f is a ground formula of $F(2)$, $F(2) \models f$, $s \models F(2)$, and $s \models f$.

Definition 3 (Ground Formulas)**.** Let F be a quantified formula and $k \in \mathbb{N}$.

$$\text{Gr}(F, k) := \{f \mid (f \text{ is a ground formula of } F(k)) \wedge (F(k) \models f)\}$$

Example 2. $\text{Gr}([\forall p, q \in P, p = q], 2) = \{(1 = 1), (1 = 2), (2 = 1), (2 = 2)\}$
Note: we sometimes use square braces to wrap formulas when it looks better than parentheses.
Notice that $\text{Gr}([\forall p, q \in P, p = q], 2)$ contains elements that are false. This indicates that the statement $[\forall p, q \in P, p = q](2)$ is not valid.

Example 3. Let sv be a state variable, then:

$$\text{Gr}((\forall p, q \in P, p \neq q \to \text{sv[p]} \neq \text{sv[q]}), 3) = \{(1 \neq 1 \to \text{sv}[1] \neq \text{sv}[1]), (1 \neq 2 \to \text{sv}[1] \neq \text{sv}[2]), ...\}$$

3 Helper Lemmas

Lemma 1. Let $k \in \mathbb{N}$, $s \in \text{States}(k)$, and F be a universally quantified formula. Then:

$$(s \models F(k)) \leftrightarrow (\forall f \in \text{Gr}(F, k), s \models f)$$

Proof. Suppose that $s \models F(k)$. For an arbitrary formula $f \in \text{Gr}(F, k)$, $F(k) \models f$ and hence we see that $s \to F(k) \wedge F(k) \to f$. It follows that $s \models f$.

Now suppose that $\forall f \in \text{Gr}(F, k), s \models f$. Suppose, for the sake of contradiction, that $s \not\models F(k)$. Then it must be the case that $s \wedge \neg F(k)$. We know that F is unversally quantified, so let $F(k) := \forall \hat{x}, \phi(\hat{x})$ where $\hat{x}$ is the vector of variables that we quantify over. Then it must be the case that $\exists \hat{x}, \neg\phi(\hat{x})$, but $\phi(\hat{x}) \in \text{Gr}(F, k)$. However this contradicts our original assumption, and hence the result is proved. □

Lemma 2. Let $k \in \mathbb{N}$, and $s \in \text{States}(k)$ such that $s \models \Phi(k)$. Then for $j \leq k$, it is also the case that $s \models \Phi(j)$.

Proof. Let k and $j \leq k$ be given and suppose that $s \models \Phi(k)$. By Lemma 1, $\forall f \in \text{Gr}(F, k), s \models \Phi(k)$. Now observe that $\text{Gr}(\Phi, j) \subseteq \text{Gr}(\Phi, k)$ due to the fact that Φ is a universally quantified PNF formula. Thus it is also the case that $\forall f \in \text{Gr}(F, j), s \models \Phi(j)$, and then the result follows from Lemma 1. □

Lemma 3. Let s be a state, f be a ground formula, and π be a permutation. Then:

$$(s \models f) \leftrightarrow (\pi(s) \models \pi(f))$$

Proof. Suppose that $s \models f$, and hence $s \to f$. By Assumption 1, $s \leftrightarrow \pi(s)$ and $f \leftrightarrow \pi(f)$, and the result follows immediately.

Now suppose that $\pi(s) \models \pi(f)$. π is a bijection–and hence invertible–thus π^{-1} is a permutation as well. By Assumption 1, $\pi(s) \leftrightarrow \pi^{-1}(\pi(s)) = s$ and $\pi(f) \leftrightarrow \pi^{-1}(\pi(f)) = f$. The result follows immediately. □

Lemma 4. Let $k \in \mathbb{N}$ and s be a state such that $s \models \Phi(k)$. If π is a permutation then it is also the case that $\pi(s) \models \Phi(k)$.

Proof. Suppose that $s \models \Phi(k)$. Then by Lemma 1, $\forall f \in \mathrm{Gr}(\Phi, k), s \models f$. But Assumption 1 shows that $s \leftrightarrow \pi(s)$ and hence $\forall f \in \mathrm{Gr}(\Phi, k), \pi(s) \models f$ which gives us our result by Lemma 1. □

4 The M-N Theorem

Theorem 1 (M-N)**.** Suppose that Φ is in PNF with only universal quantifiers, while Δ is in PNF with only existential quantifiers. Let m be the number of variables that Φ quantifies over and n be the number of variables that Δ quantifies over. If $\Phi(m+n)$ is an inductive invariant, then $\Phi(k)$ is also an inductive invariant for any $k > m+n$.

Proof. Let $k > m+n$ be given and assume that $[\Phi \wedge \Delta \to \Phi'](m+n)$ is valid. Let $s \in \mathrm{States}(k)$ such that $s \models \Phi(k)$, and let δ be a single arbitrary transition such that $\delta \models \Delta(k)$. Finally, let $f' \in \mathrm{Gr}(\Phi', k)$ be arbitrary, then, by Lemma 1, it suffices to show that $(s \wedge \delta) \models f'$.

Next, we will construct a permutation π as follows: let $x_1, ..., x_j$ be the distinct elements of P used in δ and f'. We know that $j \leq m+n$ because Δ quantifies over n variables while Φ quantifies over m variables. Let:

$$\pi := \begin{pmatrix} x_1 & x_2 & ... & x_j \\ 1 & 2 & ... & j \end{pmatrix}$$

Notice that $\pi(\delta)$ and $\pi(s)$ now only contain the elements $1, ..., j$ and, in particular, $\pi(\delta) \models \Delta(m+n)$ and $\pi(f') \in \mathrm{Gr}(\Phi', m+n)$, i.e. $\pi(f') \models \Phi'(m+n)$. By Lemma 2 we see that $s \models \Phi(m+n)$, and furthermore $\pi(s) \models \Phi(m+n)$ by Lemma 4. Thus $\pi(s \wedge \delta) \models [\Phi \wedge \Delta](m+n)$ which implies $\pi(s \wedge \delta) \models \Phi'(m+n)$ by our initial assumption. In particular, $\pi(s \wedge \delta) \models \pi(f')$ by Lemma 1, and therefore $s \wedge \delta \models f'$ by Lemma 3. □

5 Case Studies

In this section we visit several (more coming soon) distributed protocols that are parameterized by a single sort and satisfy Assumption 1.

5.1 Peterson's Mutex Protocol

Peterson's Mutex Protocol can be encoded with a transition function Δ in PNF that quantifies over two variables. A sample inductive invariant candidate is given in [1] that quantifies of two variables and works for $|P| = 2$:

```
Phi == \A p,q \in ProcSet :
   /\ pc[p] \in {"a3","a4","cs"} => flag[p]
   /\ (p#q /\ pc[p] = "cs" /\ pc[q] = "a4") => turn = p
   /\ (p # q) => ~(pc[p] = "cs" /\ pc[q] = "cs")
```

However, by the M-N Theorem, we must show that Φ is an inductive invariant for the cases when $|P| = 1, ..., 4$. In fact, we easily see that Φ fails to be inductive in the case:

```
/\ turn = 1              /\ turn = 2
/\ pc[1] = "cs"          /\ pc[1] = "cs"
/\ pc[2] = "a4"    ->    /\ pc[2] = "a4"
/\ pc[3] = "a3"          /\ pc[3] = "a4"
/\ a3(3,2)
```

References

[1] Parametric Peterson's Mutex Protocol. `https://github.com/iandardik/iinf/blob/master/ii_cutoff/mn_thm/PetersonParametric.tla`, 2022.

[2] Aman Goel and Karem Sakallah. On Symmetry and Quantification: A New Approach to Verify Distributed Protocols. In *NASA Formal Methods Symposium*, pages 131–150. Springer, 2021.