

M-N Without Permutations

Ian Dardik

April 1, 2022

1 Introduction

Finding an inductive invariant is key for proving the correctness of a distributed protocol with respect to a safety property. As such, a considerable amount of effort has been dedicated to finding and proving an inductive invariant for a given system. For example, Ivy will guide a user to interactively find an inductive invariant within the confines of a decidable fragment of FOL. In the past few years there has also been a host of research into inductive invariant synthesis for parameterized distributed protocols. The synthesis tools that remain within the bounds of a decidable logic fragment are able to guarantee that they produce an inductive invariant, however, any tool that produces a candidate inductive invariant for a system that falls outside of a decidable fragment offers no guarantee that the candidate is indeed correct. In this note, we assume that a candidate inductive invariant is *given* and we exclusively focus on the verification step.

We have discovered a syntactic class of protocols that lie outside of a decidable logic fragment, but exhibit a *cutoff* for the number of finite protocol instances which need to be verified. We have captured this result in the M-N Theorem.

In this note we begin by introducing the Sort-Quantifiers Restricted to Prenex Normal Form Language (SRPL), the logic language that we use to encode our class of protocols. We then introduce our encoding of protocols as a transition system in SRPL. Next, we will prove some key lemmas before finally presenting and proving the M-N Theorem.

2 Sort-Quantifiers Restricted to PNF Language

In this section we will define $\text{SRPL}(E, G)$ as a grammar parameterized by the sort E and the *input grammar* G .

Definition 1. Let \mathcal{V} be a countable set of variables, E be an infinitely countable sort of indistinguishable elements, and G be an input grammar that may not refer to E . A $\text{SRPL}(E, G)$ formula is defined by the grammar for the production rule of *srpl*:

arg	$::= x$	for any $x \in \mathcal{V}$
arg_list	$::= arg$	
arg_list	$::= arg, arg_list$	
Q	$::= \forall \mid \exists$	
$srpl$	$::= Q x \in E, G(arg_list)$	for any $x \in \mathcal{V}$
$srpl$	$::= Q x \in E, srpl$	for any $x \in \mathcal{V}$

The input grammar G has a single requirement—that it cannot explicitly refer to E —and therefore is quite general. We now provide an example of an input grammar to illustrate a potential use case.

Example 1. Let \mathcal{S} be a finite set of state variables, \mathcal{A} be a countable set of constants, and let \mathcal{V} be a countable set of variables. We define the grammar *sample* that is parameterized on the variable symbols x_1, \dots, x_n by the following production rules:

$$\begin{array}{ll}
\text{prim}(x_1, \dots, x_n) & ::= v & \text{for any } v \in \mathcal{S} \\
\text{prim}(x_1, \dots, x_n) & ::= y & \text{for any } y \in \mathcal{V} \\
\text{prim}(x_1, \dots, x_n) & ::= a & \text{for any } a \in \mathcal{A} \\
\text{prim}(x_1, \dots, x_n) & ::= x_i & \text{for any } 1 \leq i \leq n \\
\text{prim}(x_1, \dots, x_n) & ::= \text{prim}(x_1, \dots, x_n)[\text{prim}(x_1, \dots, x_n)] \\
\text{sample}(x_1, \dots, x_n) & ::= \text{prim}(x_1, \dots, x_n) = \text{prim}(x_1, \dots, x_n) \\
\text{sample}(x_1, \dots, x_n) & ::= \neg \text{sample}(x_1, \dots, x_n) \\
\text{sample}(x_1, \dots, x_n) & ::= \text{sample}(x_1, \dots, x_n) \wedge \text{sample}(x_1, \dots, x_n) \\
\text{sample}(x_1, \dots, x_n) & ::= \forall x \in \text{sample}(\text{arg_list}(x_1, \dots, x_n)), \text{sample}(x_1, \dots, x_n) & \text{for any } x \in \mathcal{V}
\end{array}$$

Notice that *sample* formulas have no way to refer to the sort E directly, and hence cannot quantify over E nor take its cardinality. We will use $\forall, \exists, \rightarrow$, etc. as syntactic sugar in *sample* formulas, defined in the expected way.

Definition 2 (Instance). Let E be a sort, G be a valid SRPL input grammar, ψ be a $\text{SRPL}(E, G)$ formula and let $H \subseteq E$ such that $H \neq \emptyset$. Then we define $\psi(E \mapsto H)$ by the following rules on the $\text{SRPL}(E, G)$ grammar:

$$\begin{array}{ll}
x(E \mapsto H) & ::= x & \text{for any } x \in \mathcal{V} \\
[\text{arg}, \text{arg_list}](E \mapsto H) & ::= \text{arg}, \text{arg_list} \\
[Q x \in E, G(\text{arg_list})](E \mapsto H) & ::= Q x \in H, G(\text{arg_list}) & \text{for any } x \in \mathcal{V} \\
[Q x \in E, \text{srpl}](E \mapsto H) & ::= Q x \in H, [\text{srpl}(E \mapsto H)] & \text{for any } x \in \mathcal{V}
\end{array}$$

In other words, $\psi(E \mapsto H)$ is the formula ψ with E replaced with H . We call $\psi(E \mapsto H)$ an *instance* of ψ , and when H is finite, we call $\psi(E \mapsto H)$ a *finite instance* of ψ .

Definition 3 (Finite Instance Notation). We may use a special shorthand for finite instances that mirrors the notation described in [1]. Let ψ be a $\text{SRPL}(E, G)$ formula and $k > 0$ be given. Then $\psi(k) := \psi(E \mapsto \{e_1, \dots, e_k\})$ where each $e_i \in E$ is arbitrary and distinct. We can also write $E(k) := \{e_1, \dots, e_k\}$ where each $e_i \in E$ is arbitrary and distinct.

Definition 4 (Valid SRPL Formula). Let E be a sort, G be a valid SRPL input grammar, and ψ be a $\text{SRPL}(E, G)$ formula. Then ψ is valid iff $\psi(E \mapsto H)$ is valid for every $H \subseteq E$.

Lemma 1. Let ψ be a SRPL formula. Then ψ is valid iff $\psi(k)$ is valid for all $k > 0$.

Proof. Coming soon. □

3 E -Ground Formulas

Definition 5 (ToEGround). Let E be a sort, G be a valid SRPL input grammar, and ψ be a $\text{SRPL}(E, G)$ formula. Next, let $R \subseteq \mathcal{V}$ be the variables that occur in ψ that quantify over E , let

$H \subseteq E$ such that $H \neq \emptyset$, and let $\rho : R \rightarrow H$ be given. Then we define $\text{ToEGround}(\psi, \rho)$ by the following rules on the $\text{SRPL}(E, G)$ grammar:

$$\begin{array}{lll}
\text{ToEGround}(x, \rho) & := \rho(x) & \text{for any } x \in R \\
\text{ToEGround}([arg, arg_list], \rho) & := \text{ToEGround}(arg, \rho), \text{ToEGround}(arg_list, \rho) \\
\text{ToEGround}([Q x \in E, G(arg_list)], \rho) & := G(\text{ToEGround}(arg_list, \rho)) & \text{for any } x \in \mathcal{V} \\
\text{ToEGround}([Q x \in E, srpl], \rho) & := \text{ToEGround}(srpl, \rho) & \text{for any } x \in \mathcal{V}
\end{array}$$

For this to work, we assume that each quantifier for E in ψ gets a unique variable name. This assumption comes without loss of generality since we can always alpha-rewrite duplicate quantifier variables.

Definition 6 (EGround). A formula g is an *E-ground formula* iff there exists a SRPL formula ψ and a mapping ρ such that $g = \text{ToEGround}(\psi, \rho)$. Moreover, we call g a *ground instance* of ψ .

Notice that *E-ground* formulas are not necessarily vanilla ground formulas, that is, formulas without quantifiers. We illustrate this in the following example.

Example 2. Consider the following $\text{SRPL}(E, sample)$ formula:

$$\psi := \forall x \in E, A[x] \rightarrow (\exists y \in B[x], y = 0)$$

where $A \in (E \rightarrow \{true, false\})$ and $B \in (E \rightarrow \mathcal{P}(\mathbb{N}))$ are state variables, and \mathcal{P} denotes the power set. Let $H = \{e_1, e_2, e_3\}$ and $\rho(x) = e_1$, then:

$$\text{ToEGround}(\psi, \rho) = A[e_1] \rightarrow (\exists y \in B[e_1], y = 0)$$

is an *E-ground* formula. However, it is not a ground formula because it contains a quantifier.

Definition 7 (EGr). Let ψ be a SRPL formula and let $H \subseteq E$ be finite. Then:

$$\text{EGr}(\psi, H) := \{g \mid \exists \rho, g = \text{ToEGround}(\psi, \rho)\}$$

$\text{EGr}(\psi, H)$ is the set of all possible *E-ground* formulas of the finite instance $\psi(E \mapsto H)$.

Example 3. Recall the $\text{SRPL}(E, sample)$ formula from the previous example:

$$\psi := \forall x \in E, A[x] \rightarrow (\exists y \in B[x], y = 0)$$

Let $H = \{e_1, e_2, e_3\}$, then:

$$\begin{aligned}
\text{EGr}(\psi, H) = \{ & A[e_1] \rightarrow (\exists y \in B[e_1], y = 0), \\
& A[e_2] \rightarrow (\exists y \in B[e_2], y = 0), \\
& A[e_3] \rightarrow (\exists y \in B[e_3], y = 0) \}
\end{aligned}$$

4 Transition System

Let a sort E be given along with a valid SRPL input grammar G . We encode a protocol as a transition system $T = (I, \Delta)$ where I is the initial constraint and Δ is the transition relation, both formulas encoded in $\text{SRPL}(E, G)$. We assume that I is restricted to universal quantification over E while Δ is restricted to existential quantification over E . Further assume that an inductive invariant candidate Φ is given in $\text{SRPL}(E, G)$ and is restricted to universal quantification over E . We use the notation $T(E \mapsto H) := (I(E \mapsto H), \Delta(E \mapsto H))$ where $H \subseteq E$ to denote an *instance* of T .

For the remainder of this note we will refer to E, T, I, Δ , and Φ as defined above.

Definition 8 (States).

$$\text{States}(H) := \{s \mid s \text{ is a state of } T(E \mapsto H)\}$$

In this note we consider a “state” $s \in \text{States}(H)$ to be a ground formula. More specifically, s is a conjunction of constraints that describe a single state in $T(E \mapsto H)$.

Definition 9 (Inductive Invariant). Φ is an inductive invariant iff $I \rightarrow \Phi$ and $\Phi \wedge \Delta \rightarrow \Phi'$ are valid formulas.

5 Helper Lemmas

Lemma 2. Let G be a valid SRPL input grammar and ψ be a $\text{SRPL}(E, G)$ formula restricted to universal quantification on E . Let $H \subseteq E$ be finite where $H \neq \emptyset$ and let $s \in \text{States}(H)$. Then:

$$(s \rightarrow \psi(E \mapsto H)) \leftrightarrow (\forall g \in \text{EGr}(\psi, H), s \rightarrow g)$$

Proof. TODO need to be cleaned up with latest notation.

Suppose that $s \rightarrow F(k)$. For an arbitrary formula $f \in \text{EGr}(F, k)$, $F(k) \rightarrow f$ and hence we see that $s \rightarrow F(k) \wedge F(k) \rightarrow f$. It follows that $s \rightarrow f$.

Now suppose that $\forall f \in \text{EGr}(F, k), s \rightarrow f$. Suppose, for the sake of contradiction, that $\neg(s \rightarrow F(k))$. Then it must be the case that $s \wedge \neg F(k)$. We know that F is universally quantified, so let $F(k) := \forall x_1, \dots, x_m \in P, \phi(x_1, \dots, x_m)$ where $m \geq 1$. Then, because $\neg F(k)$ holds, it must be the case that $\exists x_1, \dots, x_m \in P, \neg \phi(x_1, \dots, x_m)$. However, $\phi(x_1, \dots, x_m) \in \text{EGr}(F, k)$ which, by our original assumption, implies $\neg s$. Hence we have both s and $\neg s$ and we have reached a contradiction. \square

Lemma 3. Let G be a valid SRPL input grammar and ψ be a $\text{SRPL}(E, G)$ formula restricted to universal quantification on E . Let $H_1 \subseteq E$ where $H_1 \neq \emptyset$ and let $s \in \text{States}(H_1)$. Let $H_2 \subseteq H_1$ where $H_2 \neq \emptyset$. Then:

$$(s \rightarrow \psi(E \mapsto H_1)) \rightarrow (s \rightarrow \psi(E \mapsto H_2))$$

Proof. Suppose that $s \rightarrow \psi(E \mapsto H_1)$, it suffices to show that $s \rightarrow \psi(E \mapsto H_2)$. We know that $\psi(E \mapsto H_2)$ is in the form:

$$\psi = \forall x_1 \in H_2, \dots, \forall x_m \in H_2, F_G(x_1, \dots, x_m)$$

where F_G is a formula generated by the input grammar G . Then $s \rightarrow \psi(E \mapsto H_2)$ holds iff $s \rightarrow F_G(e_1, \dots, e_m)$ holds for arbitrary $e_1 \in H_2, \dots, e_m \in H_2$. However, this formula must hold by our assumptions that $H_2 \subseteq H_1$ and $s \rightarrow \psi(E \mapsto H_1)$ where ψ is universally quantified over E . \square

Lemma 4. Let G be a valid SRPL input grammar and ψ be a $\text{SRPL}(E, G)$ formula restricted to existential quantification on E . Let $H_1 \subseteq E$ where $H_1 \neq \emptyset$. Let $g \in \text{EGr}(\psi, H_1)$, and let e_1, \dots, e_m be the elements of H_1 that occur in g . Then for any $H_2 \supseteq \{e_1, \dots, e_m\}$:

$$(g \rightarrow \psi(E \mapsto H_1)) \rightarrow (g \rightarrow \psi(E \mapsto H_2))$$

Proof. Suppose that $g \rightarrow \psi(E \mapsto H_1)$, then it suffices to show that $g \rightarrow \psi(E \mapsto H_2)$. We know that ψ is of the form:

$$\psi = \exists x_1 \in H_2, \dots, \exists x_m \in H_2, F_G(x_1, \dots, x_m)$$

where F_G is a formula generated by the input grammar G . Because $g \rightarrow \psi(E \mapsto H_1)$, it must be the case that e_1, \dots, e_m witness the existential quantifiers of $\psi(E \mapsto H_1)$. However, $\{e_1, \dots, e_m\} \subseteq H_2$, and hence $g \rightarrow \psi(E \mapsto H_2)$. \square

6 The M-N Theorem

In this section, we will establish initiation and consecution in two separate lemmas. The M-N Theorem is then easily proved from these two lemmas.

Lemma 5 (M-N Initiation). Let m be the number of quantifiers over E in I . Then if $I(m) \rightarrow \Phi(m)$ is valid, $I(k) \rightarrow \Phi(k)$ is also valid for all $k > m$.

Proof. Coming soon. □

Lemma 6 (M-N Consecution). Let m be the number of quantifiers over E in Φ and n be the number of quantifiers over E in Δ . Then if $\Phi(m+n)$ is inductive, $\Phi(k)$ is also inductive for any $k > m+n$.

Proof. Assume that $[\Phi \wedge \Delta \rightarrow \Phi'](m+n)$ is valid. Let $k > m+n$ be given, we want to show that $[\Phi \wedge \Delta \rightarrow \Phi'](k)$ is also valid. Let $H = \{e_1, \dots, e_k\} \subseteq E$ be an arbitrary finite instance of E . Let $s \in \text{States}(H)$ such that $s \rightarrow \Phi(E \mapsto H)$ and let $\delta \in \text{EGr}(\Delta, H)$ such that $\delta \rightarrow \Delta(E \mapsto H)$. Then $(s \wedge \delta)$ is an E -ground formula that describes the states reachable from s in one “ δ step”, and it suffices to show that $(s \wedge \delta) \rightarrow \Phi'(E \mapsto H)$. Furthermore, let $\phi' \in \text{EGr}(\Phi', H)$ be arbitrary, then, by Lemma 2 and the fact that Φ' is restricted to universal quantification on E , it suffices to show that $(s \wedge \delta) \rightarrow \phi'$.

Let $\alpha_1, \dots, \alpha_i$ be the unique elements of $\{e_1, \dots, e_k\}$ that occur in $(\phi \wedge \delta)$, then we know that $i \leq m+n$ because $\phi \in \text{EGr}(\Phi, H)$ where Φ quantifies over m variables and $\delta \in \text{EGr}(\Delta, H)$ where Δ quantifies over n variables. Let $j = m+n-i$, then we can choose β_1, \dots, β_j such that $\{\beta_1, \dots, \beta_j\} \subseteq (\{e_1, \dots, e_k\} - \{\alpha_1, \dots, \alpha_i\})$ (define $\{\beta_1, \dots, \beta_j\} = \emptyset$ in the case where $j = 0$). Notice that $|\{\alpha_1, \dots, \alpha_i, \beta_1, \dots, \beta_j\}| = m+n$, and hence, by our initial assumption:

$$[\Phi \wedge \Delta \rightarrow \Phi'](E \mapsto \{\alpha_1, \dots, \alpha_i, \beta_1, \dots, \beta_j\})$$

must be a valid formula.

Now, $s \rightarrow \Phi(E \mapsto \{\alpha_1, \dots, \alpha_i, \beta_1, \dots, \beta_j\})$ due to Lemma 3 because Φ is restricted to universal quantification on E . Furthermore, $\delta \rightarrow \Delta(E \mapsto \{\alpha_1, \dots, \alpha_i, \beta_1, \dots, \beta_j\})$ by Lemma 4 because Δ is restricted to existential quantification on E . Thus we see:

$$(s \wedge \delta) \rightarrow [\Phi \wedge \Delta](E \mapsto \{\alpha_1, \dots, \alpha_i, \beta_1, \dots, \beta_j\}) \rightarrow \Phi'(E \mapsto \{\alpha_1, \dots, \alpha_i, \beta_1, \dots, \beta_j\}) \rightarrow \phi'$$

□

Next we present the M-N Theorem:

Theorem 1 (M-N). Let m be the number of quantifiers over E in Φ and n be the number of quantifiers over E in Δ . Then if $\Phi(m+n)$ is an inductive invariant, $\Phi(k)$ is also an inductive invariant for any $k > m+n$.

Proof. This follows immediately from Lemma 5 and Lemma 6. □

Perhaps even more important than the M-N Theorem itself, is the following corollary:

Corollary 1. Let m be the number of quantifiers over E in Φ and n be the number of quantifiers over E in Δ . Then if $\Phi(k)$ is an inductive invariant for all $k \in \{1, \dots, m+n\}$, then Φ is an inductive invariant for T .

Proof. Suppose that $\Phi(k)$ is an inductive invariant for all $k \in \{1, \dots, m+n\}$. By the M-N Theorem, we know that $\Phi(k)$ is also an inductive invariant for all $k > 0$. The result then follows from Lemma 1. (TODO: a bit more is need here) □

References

- [1] Aman Goel and Karem Sakallah. On Symmetry and Quantification: A New Approach to Verify Distributed Protocols. In *NASA Formal Methods Symposium*, pages 131–150. Springer, 2021.